

Entity Relationship Model

This model represents entities of a database i.e. helps in identifying entities to be represented in a database and also represents how those entities are related.

Rectangles → Represents strong entities in ER model
that can exist by themselves. □

Double rectangles → Represents weak entities in ER model that need to be related to a strong entity to be defined, as by themselves they have no existence. □

Ellipse → Represents attributes in ER model. ◻ ○

Double ellipse → Represents multi-valued attributes in ER model. ○○

Diamond → Represents relationships among entities. ↗
ER model

Entity

Strong entity
Weak entity

Attribute

Key attribute
Composite attribute

Multi-valued attribute

Derived attribute

Relationship

One to one

Many to many

One to many

Many to one

Strong
a stron
the com
Weak ent
are all
on the
But w
w.r.t.
Identify
entity a
relations)

Employ

Attribute

Eg: R

• Key a

• Comp

A

• Mult
than
can b

(Ph

Entity → It's an object with physical (car, house etc.) existence or conceptual (company, job etc.) existence.

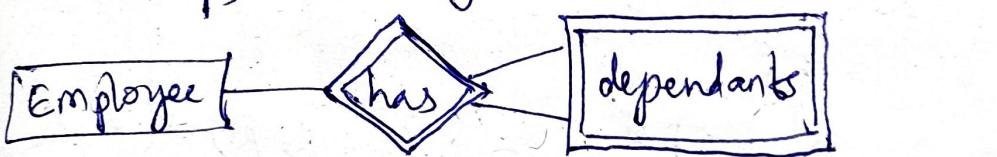
Entity set → A set of entities, all of the same type.

Strong entity → Eg An employee of a company is a strong entity as he/she is directly associated with the company.

Weak entity → Eg Employee's wife, child, parents, they are all dependants of the employee hence dependent on the company through employee.

But without the employee, they have no significance w.r.t the company, hence they won't exist.

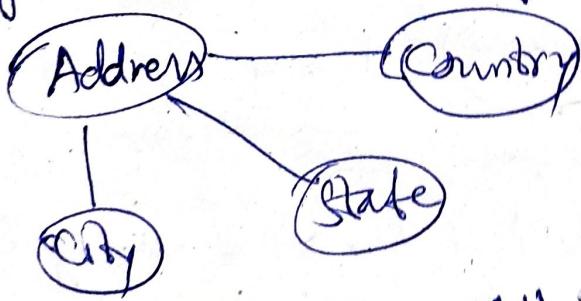
Identifying Relationship → The relationship b/w weak entity and its identifying strong entity is called identifying relationship, shown by double diamond.



Attributes → Properties that define the entity type.

Eg Roll no; DOB, address that define student.

- Key attribute → Uniquely identifies entity.
- Composite attribute → composed of other attributes.



- Multivalued attribute → Attribute consisting of more than one value for an entity. Eg a student can have multiple phone nos.

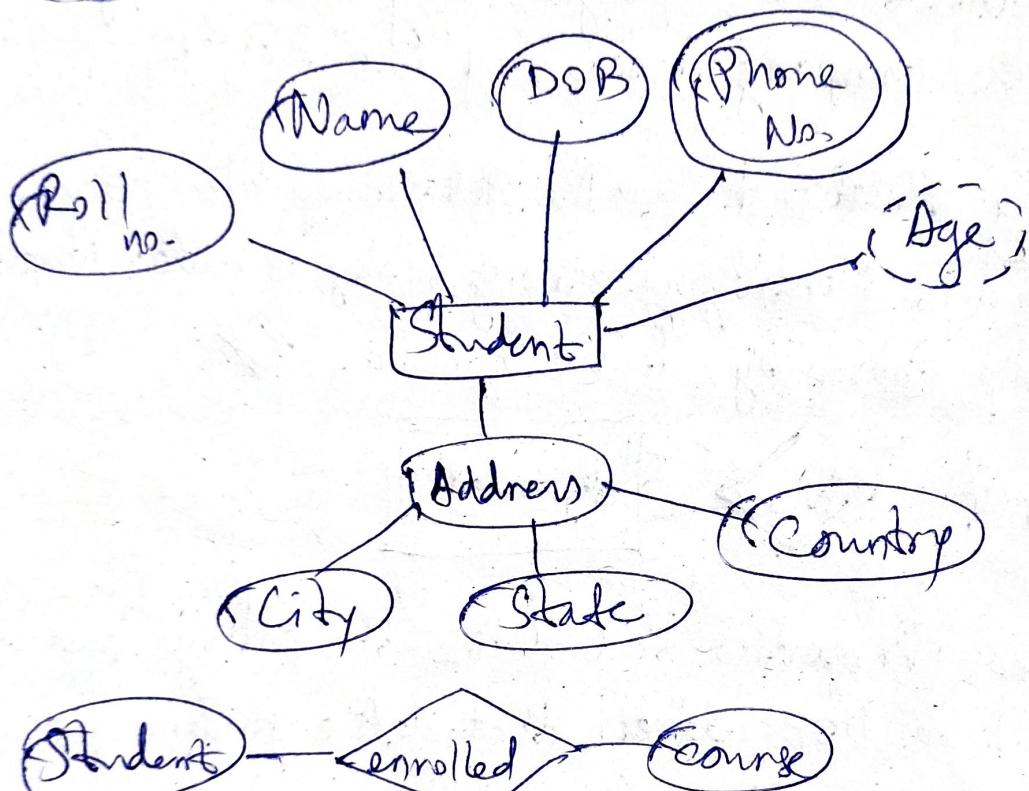


* Derived attribute +

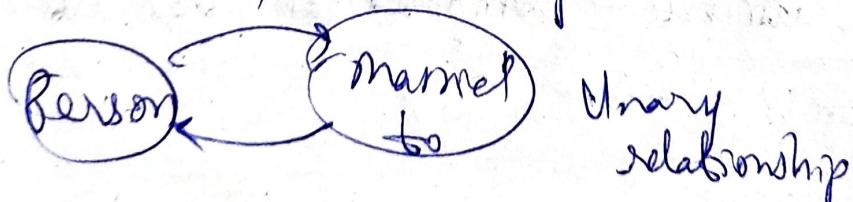
An attribute that can be derived from another attribute of an entity. E.g. DOB derived from Age
E.g. age derived using DOB

→ Dotted ellipse

(Age)

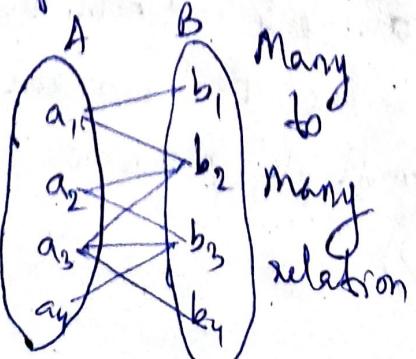
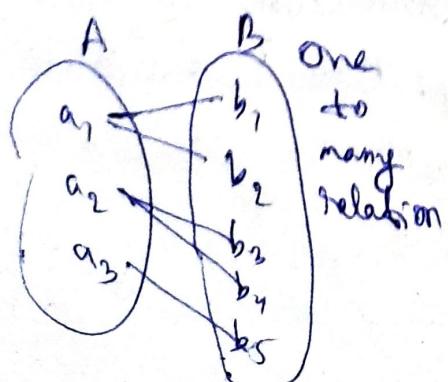


Binary relationship



Unary relationship

If no. of relationships b/w multiple entity sets is called a relationship set.



Database System Concepts

Korth & Silberschatz

Navathe

Ramkrishnan Gehrke

CJ Date

Database Management System

Data storage

Data retrieval (fetch)

Data updation / alteration

Structured data → Data present in fixed format

Data abstraction ←

4 levels of data abstraction :

① Physical



② Conceptual

Data is conceptualised in the form of a table with several columns for different parameters. ①, ② → Theory

③ View

It refers to the different levels of accessibility to a certain database, pertaining to who gets access to a certain data, and who's denied.

④ User

Final front-end view of data that's shown to the final user. ③, ④ → Lab

Database Design

① ER model (Entity Relationship model)

② Relational model

Entity → An object with physical existence, or conceptual existence.

Collection of similar entities → Entity set

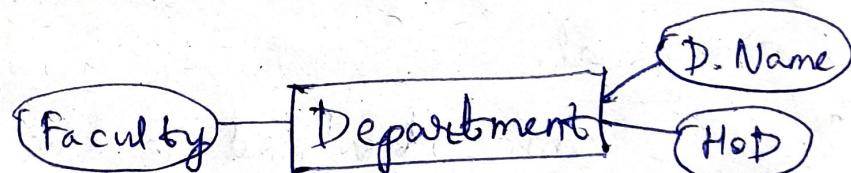
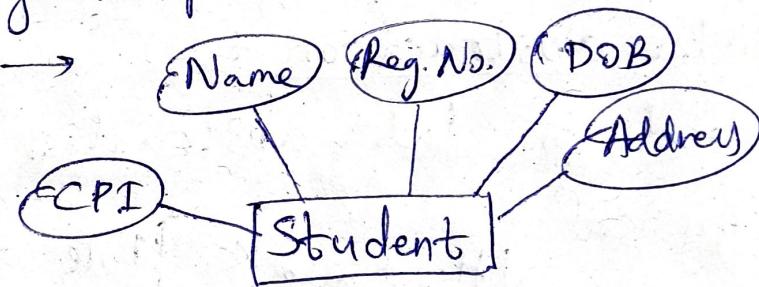
Every entity has certain distinguished features, attributes is another name for features.

Entity → \square , attribute → \circ
Rectangular shape Oval shape

Dean Academic →

↳ Student

↳ Departments



Key attributes

- | | |
|-----------------|---------------|
| ① Candidate key | ③ Foreign key |
| ② Primary key | ④ Super key |

Key attribute → Uniquely identifies an entity-

For a student, Reg. No. is the primary key as for every student, Reg. No. is unique.

There can be several other attributes associated with Reg. No. but it'll still be a key, although the rest will be redundant.

Candidate key → Set of attributes for an entity such that it uniquely identifies the entity.

Primary key → The candidate key containing the minimum no.

Superkey → Along with essential attribute, it contains redundant attributes which can be removed, but the property of being a key still remains.

DBMS

21/8/23

Domain → The range of permissible values from which its corresponding attribute can accept any one, is called the domain of that attribute.

Foreign key → A key that relates to attributes belonging to different entities.

ER model → Entity
Relational model → Tuple
Table → Row

SQL Table &
Student

{ varchar RegNo }
varchar Name }
varchar course }
int sem }

RegNo Primary
Key

Underlined in
ER model



Candidate Key's

Sets of different attributes, where ~~only~~ all attributes in every set together uniquely identifies a student i.e. a row.

Students $\{ \text{Reg. No} \}$ \leftarrow 2 such
 $\{ \text{Name, Address} \} \leftarrow$ ~~all~~ candidate
Keys

I cannot remove any ^{one} attribute from this set, as then it'll lose its property of being a key.

If I remove address, only name is insufficient for uniquely identifying a student, i.e., a row, as multiple students can have same name.

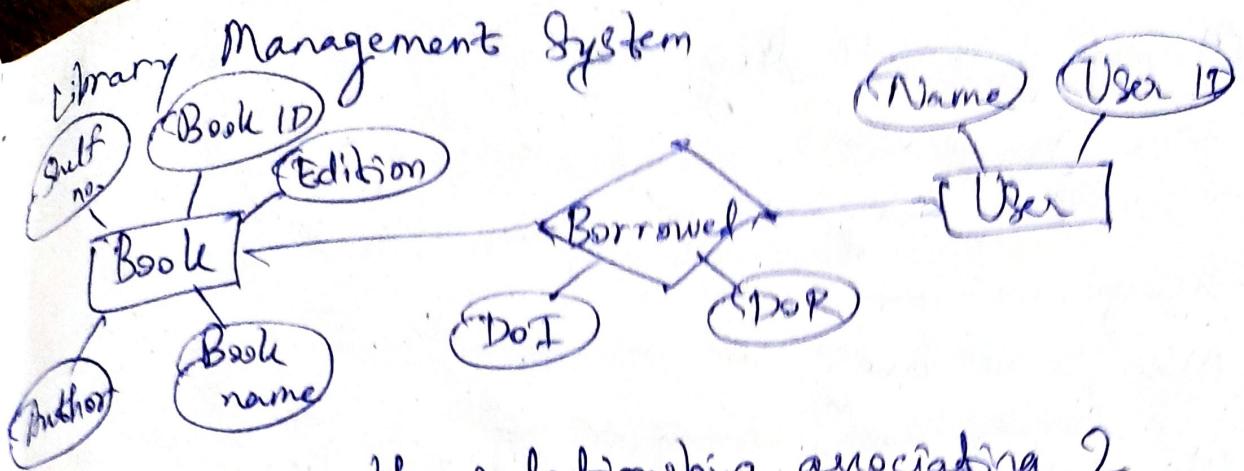
Primary key → One such set of attributes from candidate keys can be a primary key.

Usually, primary key is that set of attributes whose size is minimum.

So, here $\{ \text{Reg. No} \}$ is preferred.

Super Key → Set of attributes such that it's a key but it contains at least one or more redundant attributes, upon removal of which, it still remains a key.

Eg → $\{ \text{Reg. No, Name} \}$ is a super key as I can remove Name from it but it'll still be a key.



Borrowed is the relationship associating 2 entities Book & User in the ER model.

Now, the primary key of the relationship Borrowed is the combination of primary keys of its associative entity sets.

$$\text{Book} \rightarrow \{\text{Book ID}\} \rightarrow \{\text{Book ID, User ID}\}$$

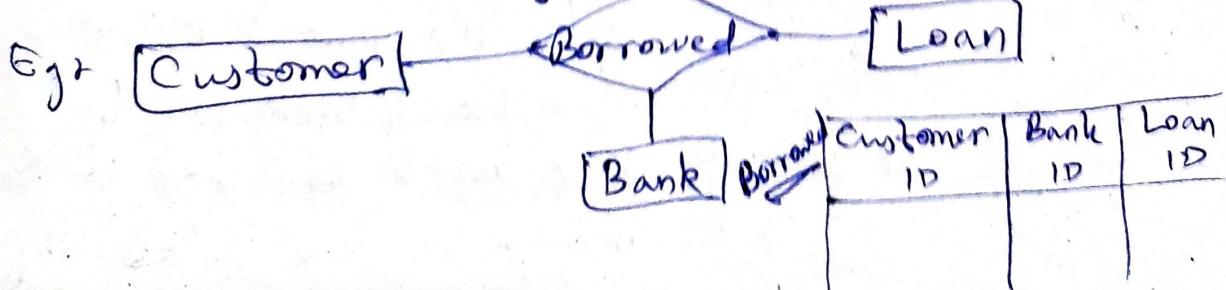
$$\text{User} \rightarrow \{\text{User ID}\}$$

And the relationship itself has its own attributes, Date of Issue(DoI) and Date of Return(DOR).

In the table format, ~~Borrowed~~ Borrowed is known as a relation, not relationship.

In table format, the 'Borrowed' relationship will have 2 columns, Book ID & User ID i.e. the primary keys of the entity sets it's associated with.

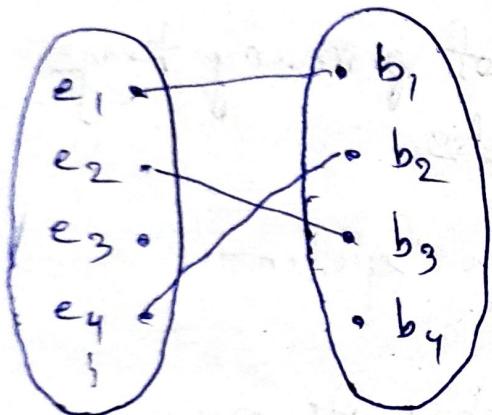
A relation can be associated with 2 or more entities simultaneously.



Mapping Cardinality

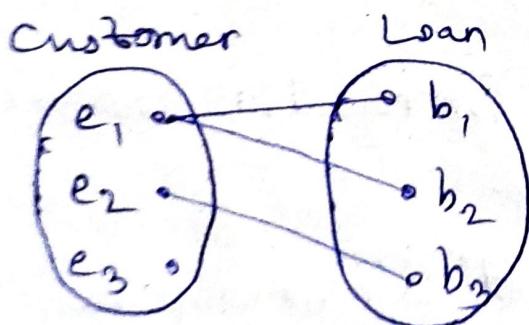
- One-to-one
- One-to-many
- Many-to-one
- Many-to-many

Mapping b/w 2 entity sets, where one or more entities of 1 set are in a relationship with one or more entities of the 2nd set.



A one-to-one mapping cardinality, it is NOT mandatory that every entity in a set has to have a relationship with atleast one entity in the 2nd set.

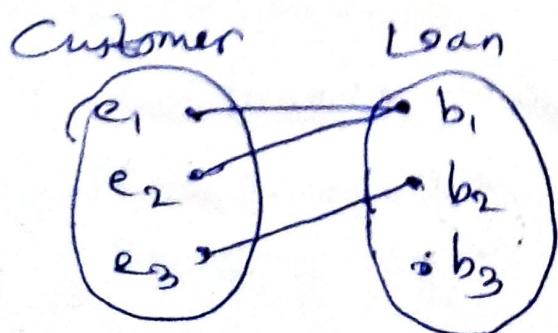
One-to-many



One customer can take several loans, but only he/she has to pay it back, and one loan ID can only be assigned to one person, so ~~many-to-one~~

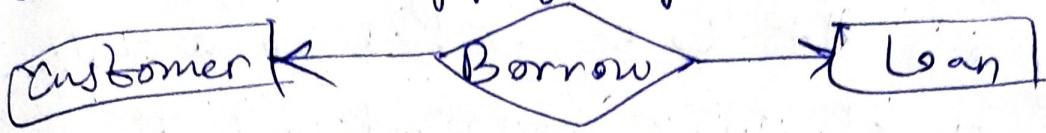
it's a one-to-many mapping.

Many-to-one



Every person can take only one loan, BUT they can have nominees to pay back the loan, so a many-to-one relationship.

One-to-one mapping depiction +



One-to-many +



Many-to-one



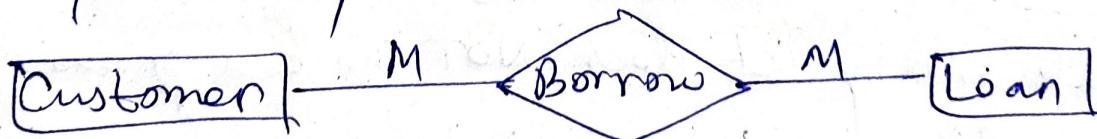
Many-to-many



One-to-many



Many-to-many

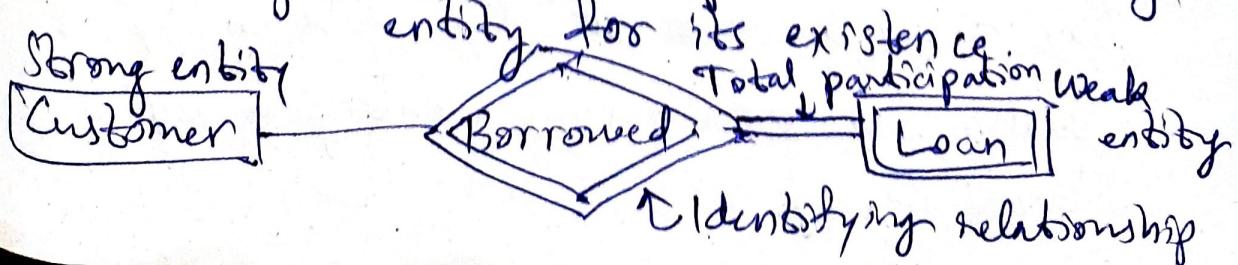


Existential Dependency

Some entities can't exist by themselves, they need to depend on another entity for their existence.

Strong entity → One that can exist on its own.

Weak entity → One that depends on a strong entity for its existence.



The King has to Empire either
keep the Indians to the North
or Europe which
would keep Europe out
of Empire but this would
cause Indians to the South
to Europe out of Empire
so Europe out of Empire
is not wanted but the
other Indians to the South
and Europe out of Empire
is not wanted
so Europe out of Empire

Total participation → Every loan ID in the entity set HAS to be related to a customer ID, so there's 100% participation in the loan entity sets to customer entity sets, hence the solid line shows total participation.

Cardinality

One - to - one

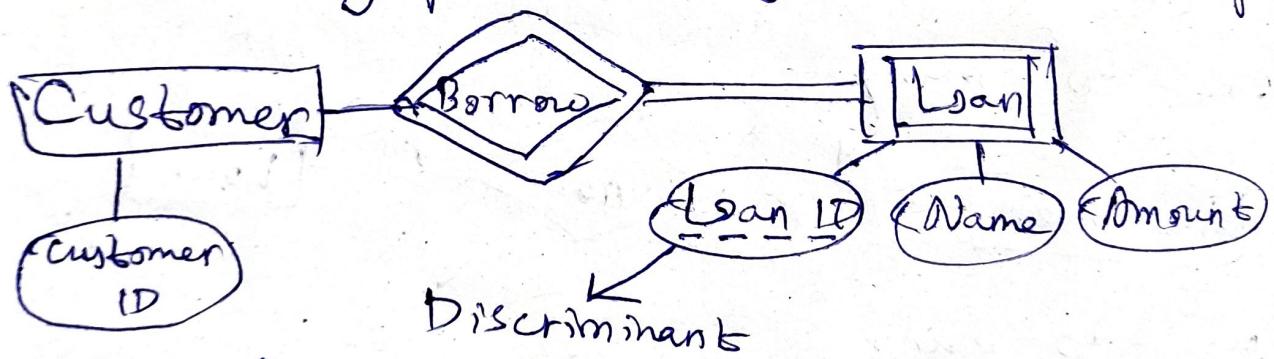
One - to - Many

Many - to - one

Many - to - Many

Mapping Cardinality →

Shows which members in each entity set is actually participating in the relationship.



Weak entity set does NOT have a primary key of its own, it needs to associate with primary key of its strong entity.

Loan			
Customer ID	Loan ID	Name	Amt

While loan ID seems like the primary key of loan, it is not. It's rather called

the discriminant of the weak entity as it seems to be a unique identifying key of that entity.



Customer → Strong entity has → identifying relationship
 Account → weak entity

Many-to-many mapping ↗

One customer can have several accounts.

And one acc. can be a joint account, linked to many customers.

DBMS

8/8/23

Constraints :

- 1) Domain constraint 3) Entity integrity constraint
- 2) Key constraint 4) Referential integrity constraint
- 1) Value of an attribute has to be confined within the domain of the attribute.
- 2) No 2 tuples in a relation i.e. no 2 rows in a table can have the same values for a primary i.e. defining key. E.g. Reg. No. can't be same

Reg. No.	Name
29	Ram
25	Ram

- 3) The defining key attribute of a tuple can NEVER be a null value.

Primary key can never be a null value.

↑	Emp. Id	Name	Dep. No.	Address	↓	
	Primary key (underlined)	X	3	Allahabad		
20	21	Y	5	Lucknow	3	Research
	22	Z	5	Kanpur	5	Admin
					DEPARTMENT	

The Dep. No. of Employee relation is a foreign key referencing to the Dep. No. of the Department relation.

Referential integrity dictates that the Dep. No. of Employee Table HAS to reference to a valid Dep. Name in the Department table.

Foreign key can be NULL, as it references to the primary key in the Department table, but it can NOT be a value that does NOT exist as a primary key in the Department table, and needless to say that the primary key can never be NULL.

Also, foreign key need not always be unique.

Employee table → Referencing relation

Department table → Referenced relation

This is the referential integrity constraint.

A set of attributes FK (foreign key) in relation schema R_1 is a foreign key of R_1 that references R_2 if it satisfies 2 conditions :-

- ① The attributes in FK have the same domain as that of PK of R_2
- ② A value of FK in a tuple t_1 either occurs as a value of PK for some tuple t_2 or is null.

In the former case, we have

$$t_1[FK] = t_2[PK]$$

$R_1 \rightarrow$ referencing relation

$R_2 \rightarrow$ referenced relation

If ① and ② satisfies, we say that referential integrity holds.

$t_i \rightarrow$ tuple (or row)

Eg: $\{20, X, 3, Allahabad\}$ is a tuple

$t_2 \rightarrow$ tuple (or row)

Eg: $\{3, Research\}$ is a tuple

1), 2), 3) are defined on single relations.

4) is defined on 2 relations.

$\langle 20, A, 3, Kanpur \rangle$ is invalid insertion, as a row with employee ID 20 (primary key) already exists, so key constraint.

$\langle X, 22, 3, Lucknow \rangle$ is invalid insertion. Firstly, we insert ordered tuple, hence X \notin domain of emp-ID, so domain constraint is violated.

Why is a table called a relation?

Relation is basically subset of Cartesian product of several sets which form many ordered tuples.

Here, imagine cross-product given below,

Emp-ID \times Name \times Dep-No. \times Address

and the relation (or table) hence formed R

$R \subseteq \text{Emp-ID} \times \text{Name} \times \text{Dep. No.} \times \text{Address}$
set of set of set of set of
nos. strings nos. strings

Order of Deletions →

There's a certain sequence of deletion of any tuple in a table i.e. you CANNOT delete a tuple whose primary key is being referenced to by another relation i.e.

if a primary key has a dependency it cannot be straightforwardly deleted until its dependencies are taken care of.

I can't straightforwardly delete $\langle 3, \text{Research} \rangle$ from Department table without deleting $\langle 20, X, 3, \text{Allahabad} \rangle$ in Employee table first. But I can delete the latter first, without any worry as it's not referenced to by any other table.

We have company database.

Company divided into departments.

Each dept. has name, number and an employee who manages Dept.

We keep track of start date of Dept. Manager.

A dept. may have several locations.

Each dept. controls a no. of projects.

Each project has a unique name, unique no. and is located at a single location.

We store each employee's SSN, address, name.

Each employee works for 1 dept. but may work on several projects together.

We keep track of the no. of hours per week that an employee currently works on, for each project.

We also keep track of direct supervisor of each employee.

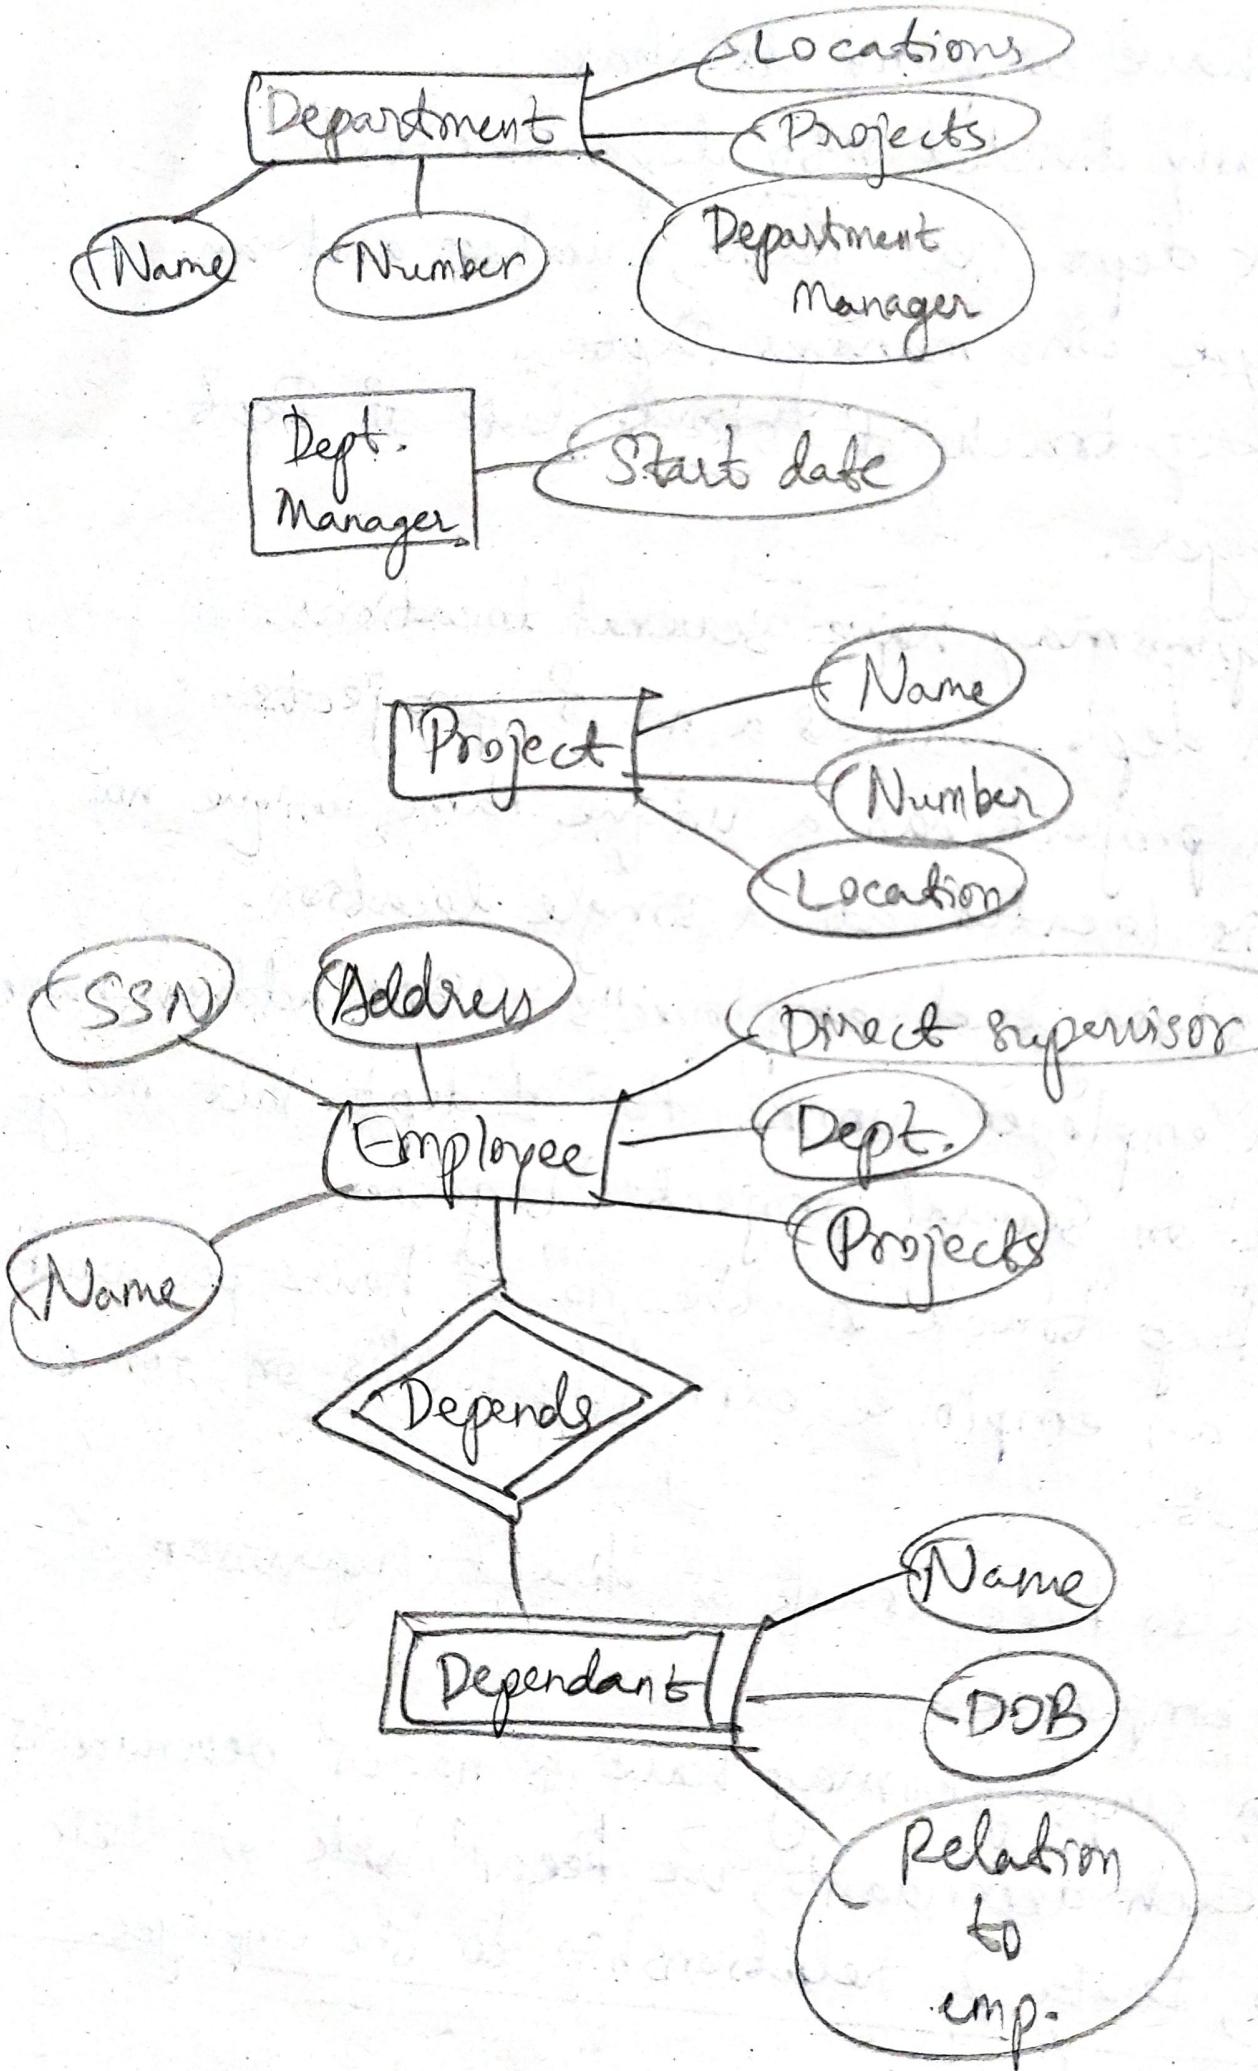
Each employee may have a no. of dependants. for each dependant, we keep track of their name, DOB & relationship to the employee.

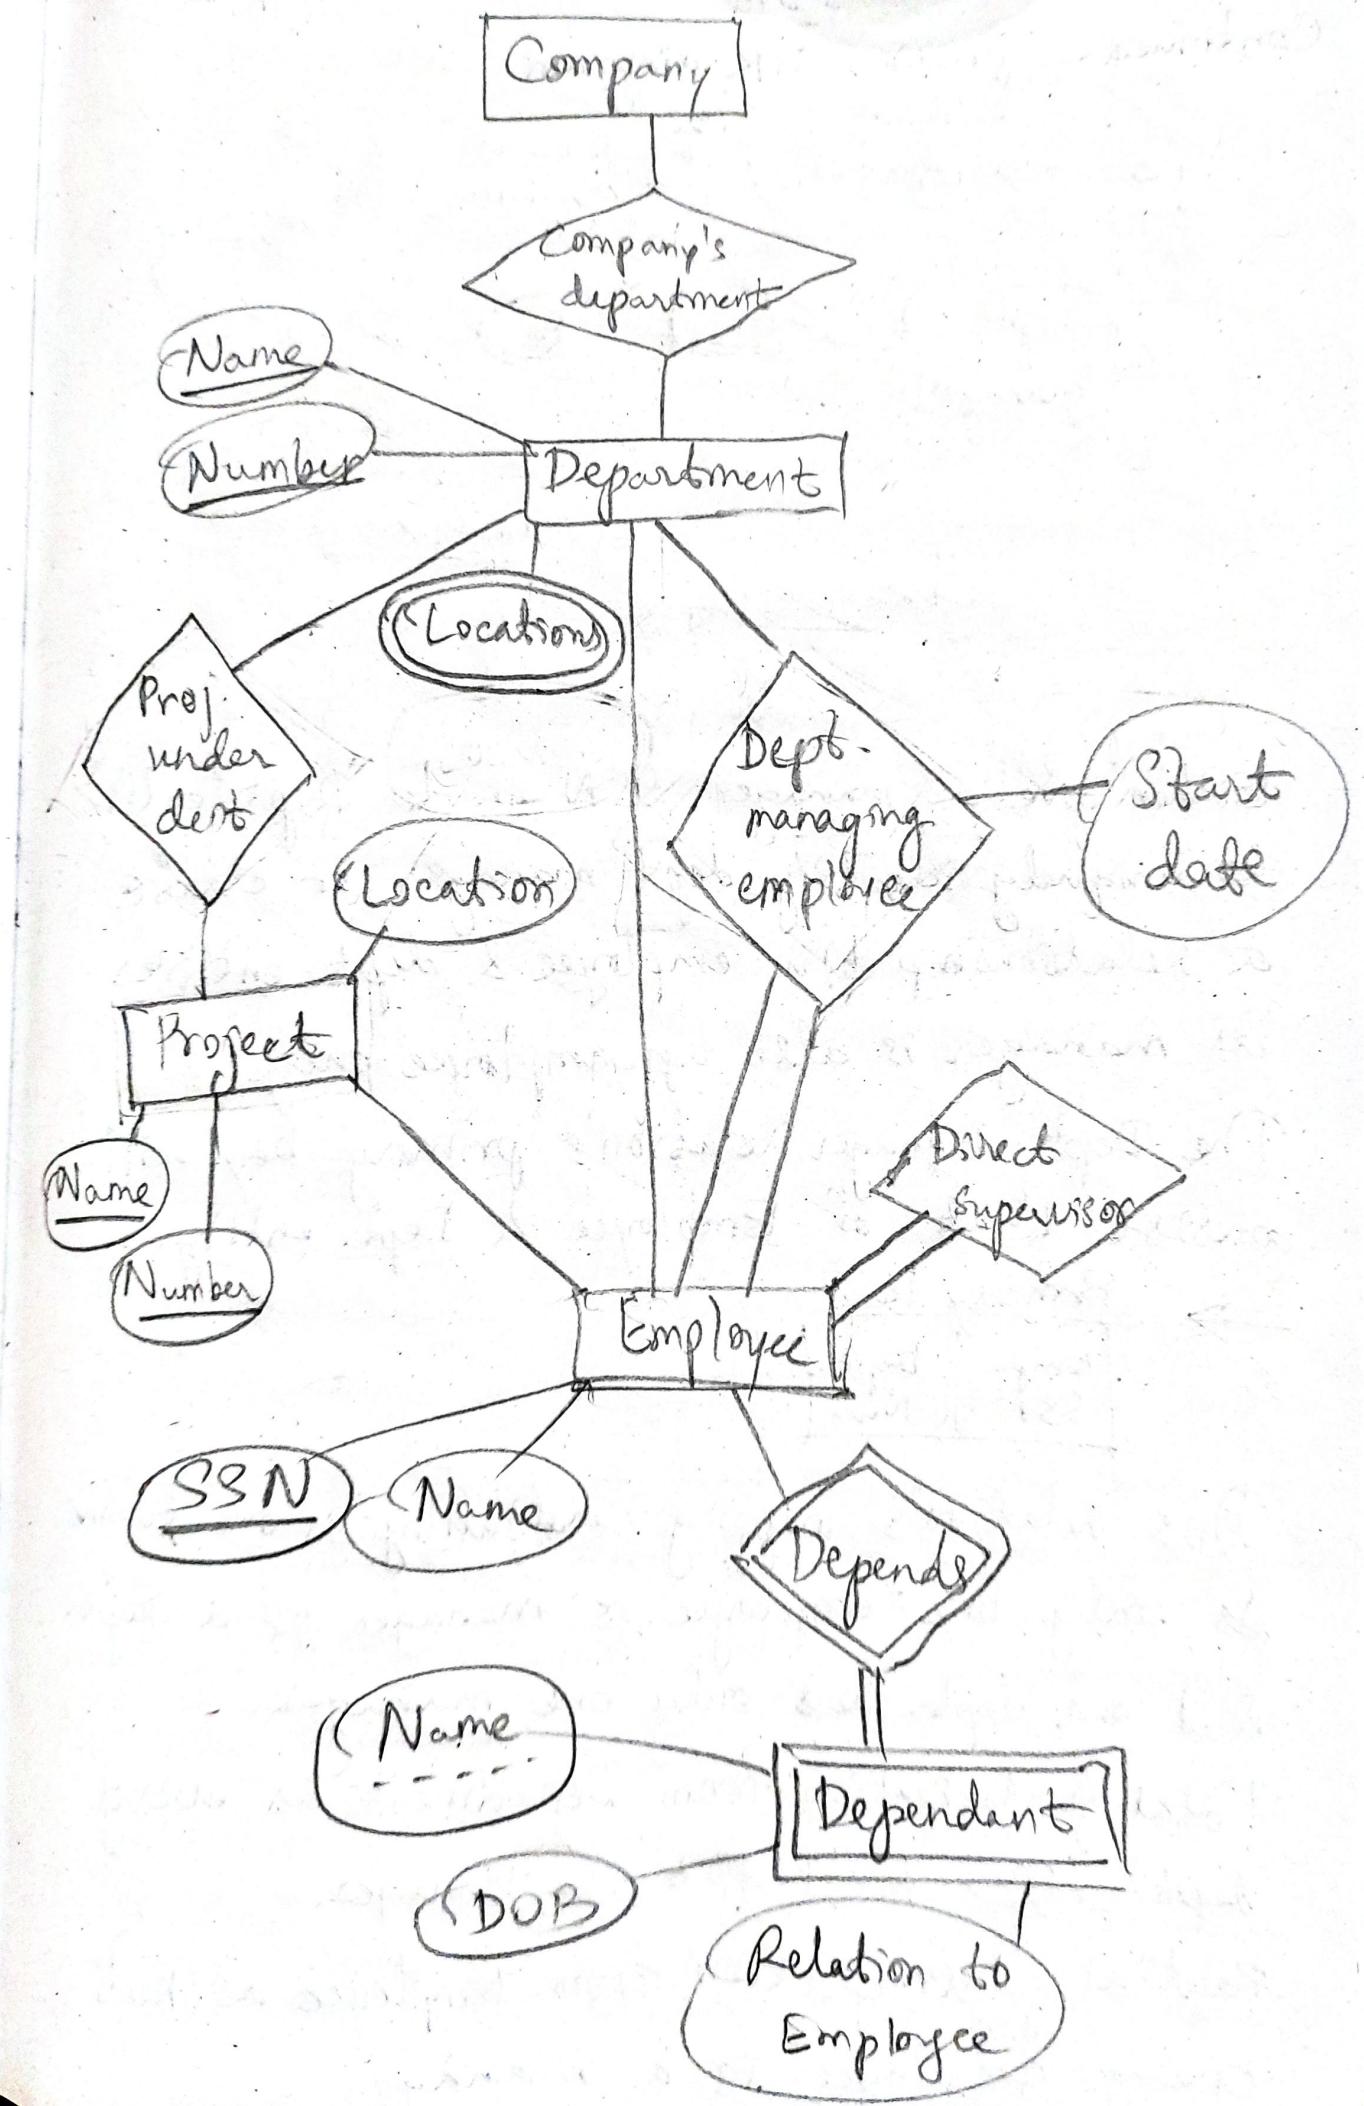
Entities → Dept. Manager

Dept.

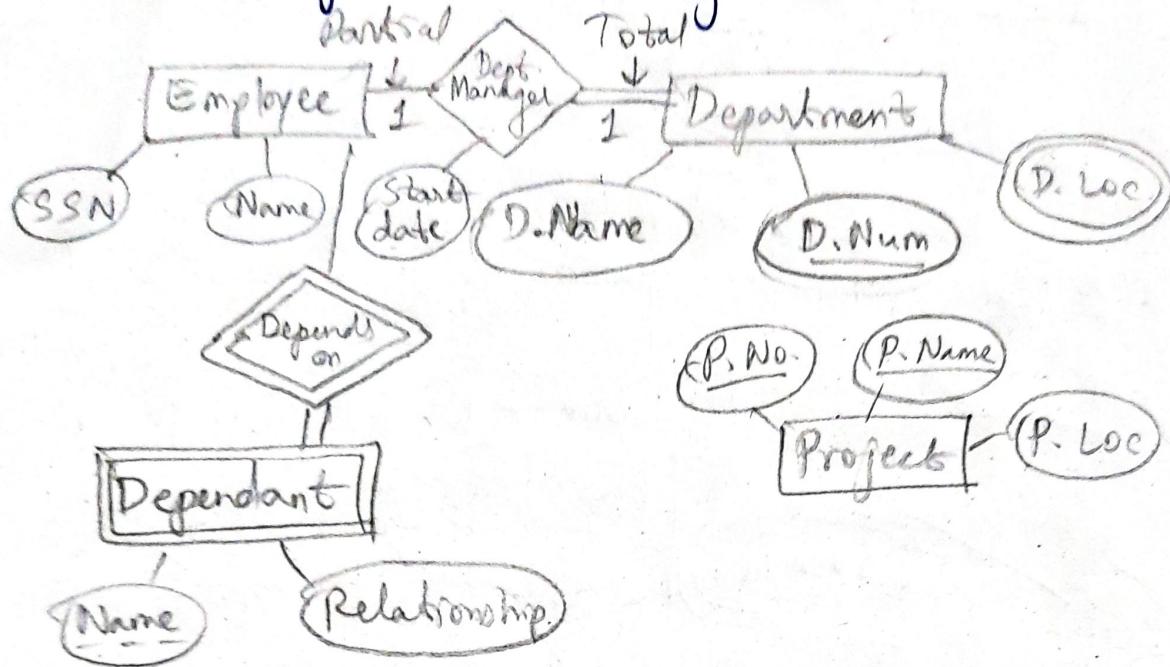
Company

Project



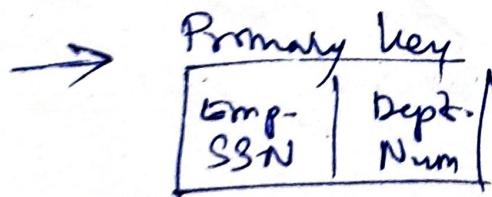


Continued... previous ER diagram



Given a manager SSN no. to Dept. entity to uniquely identify dept. manager, or create a relationship b/w employee & dept. entities as manager is also an employee only.

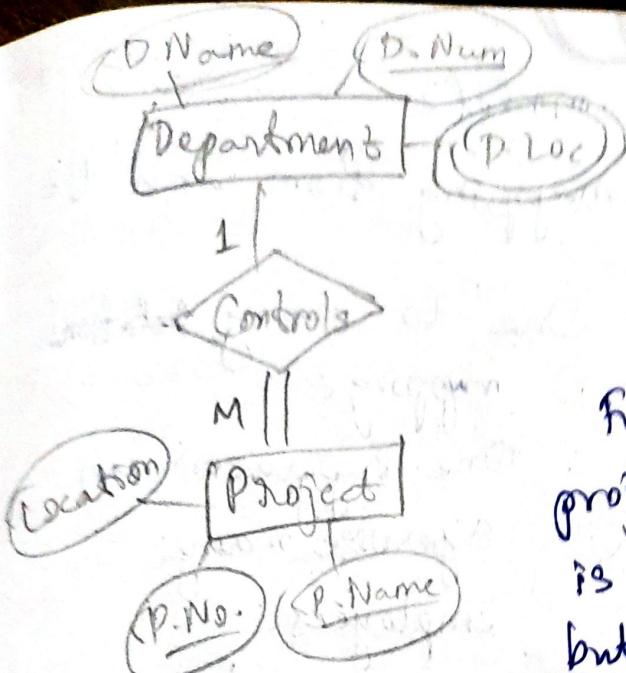
The Dept. Manager relation's primary key will consist of Employee & Dept. entity.



This relation's mapping cardinality is one-to-one. As only one employee is manager of a dept. And one dept. has only one manager.

Total participation from Department as every dept. HAS TO HAVE a manager.

Partial participation from Employee as not every employee is a manager.



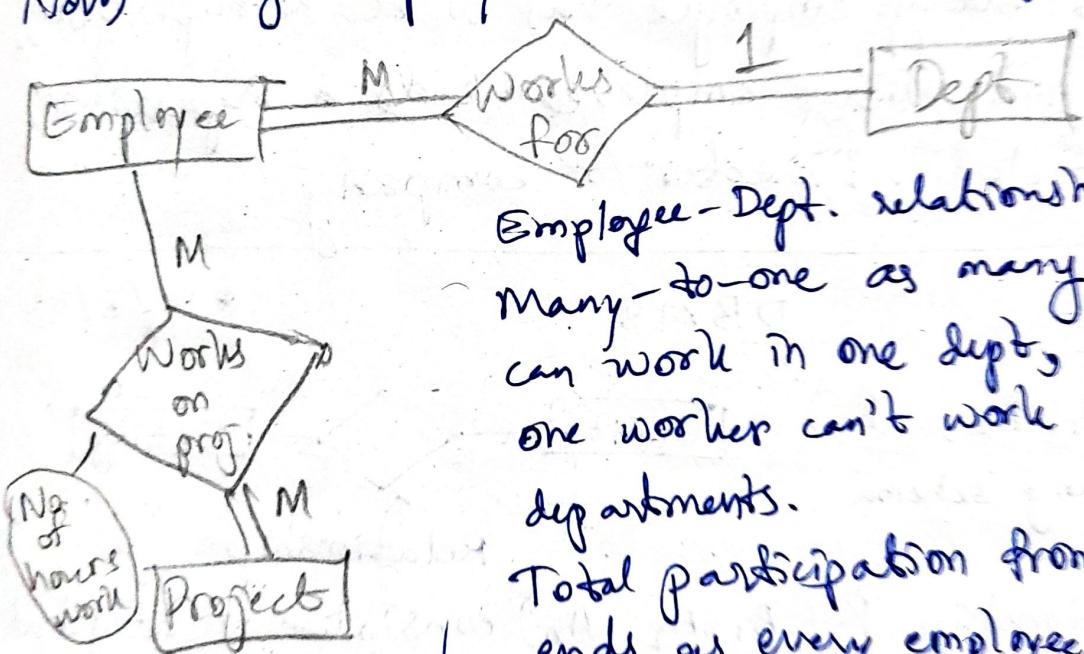
Mapping cardinality

One-to-many

as one dept. may control many projects.

Full participation from project entity as every proj. is under control of a dept. but not vice-versa. A dept. may've finished all projects.

Now, every employee works for a dept.



Employee - Dept. relationship is Many-to-one as many employees can work in one dept, but one worker can't work in many departments.

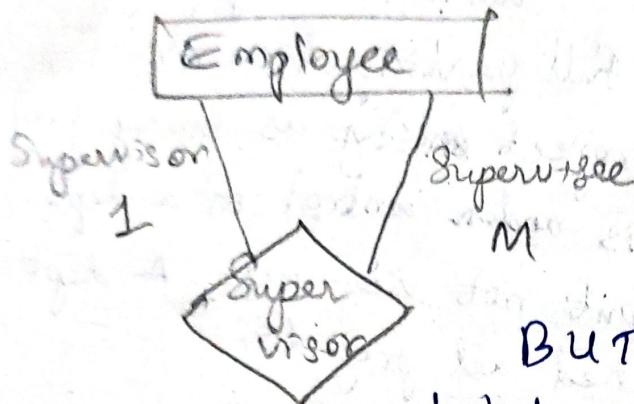
Total participation from both ends as every employee has to be part of a dept and every dept has to have an employee atleast.

Employee - Project
Many-to-many & as many employees can work on a project, and one employee can work on many projects.

Total participation from proj. → Every proj. needs to have atleast one employee work on it, but not every employee needs to work on a proj. Eg. HOD

Employee & Supervisor model

Unary relationship mapping from and to the employee.



One-to-many ~~relations~~ mapping +

One supervisor can supervise many employees.

BUT, neither end has total participation.

Not every employee has to be a supervisor, and not every employee needs a supervisor, e.g. HOD, Director of company

DBMS

• 10/8/23

$R(R_1, R_2, \dots, R_n)$

company schema



Relationship

Relationship $R_i(A_1, A_2, \dots, A_n)$ consists of attributes
Strong entity sets are converted to relations.

Employee

SSN	Name	Address

Department

D. No.	Dep. Name

Project

P. No.	P. Name	P. Loc.

Dependant

SSN	Name	Relationship

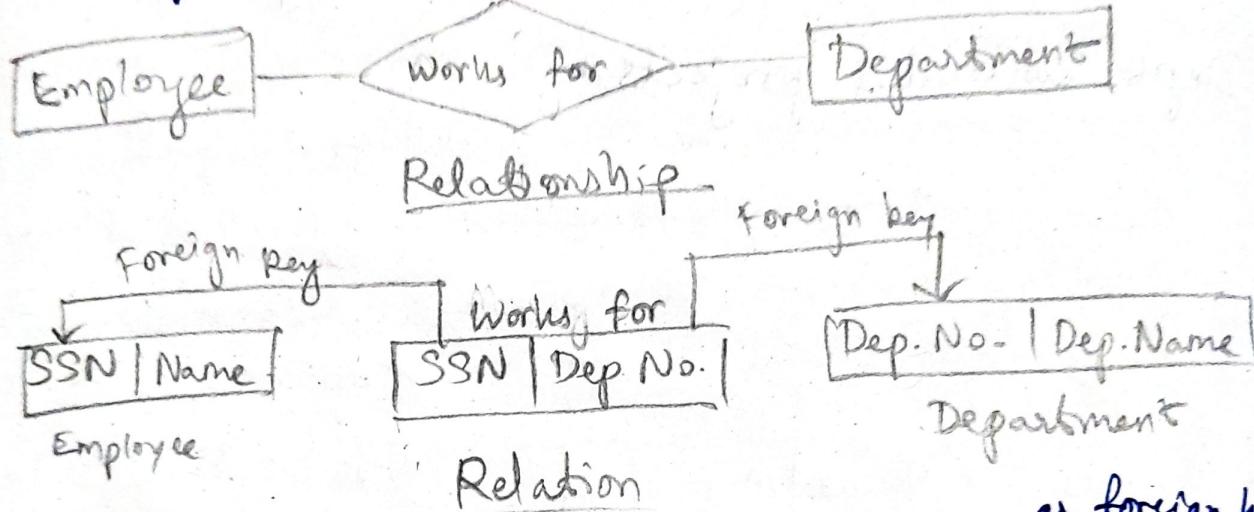
Primary key
of strong entity Discriminant

Above 4 entities are shown as relations. Now, a multi-valued attribute will have to be shown as a separate relation like department locations.

Dept. Locations

D. No.	D. Loc.

In a row, one column can accept only one entry. So, for a multi-valued attribute it'll take multiple columns, especially when you don't know the no. of columns you'll need. So, better you create a separate row for that and {dept. no., location} together forms a key.



Alternative → You can add an attribute Dep. No. to Employee entity or add SSN of employee as foreign key to Department entity.

The former is a better idea as one employee works in only 1 department, but one department has many employees working under it.

Here, Dep. No. is foreign key of employee entity that ~~maps to~~ references to Department entity.

Employee

SSN	Name	Address

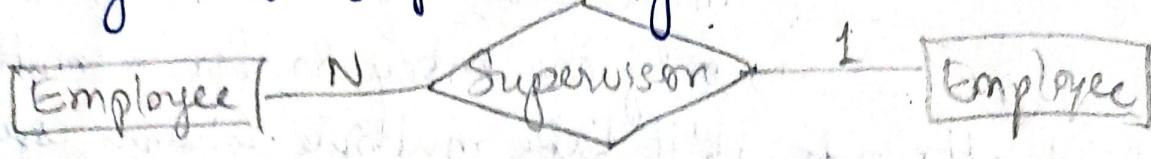
Department

D. No.	D. Name	Mgr. SSN

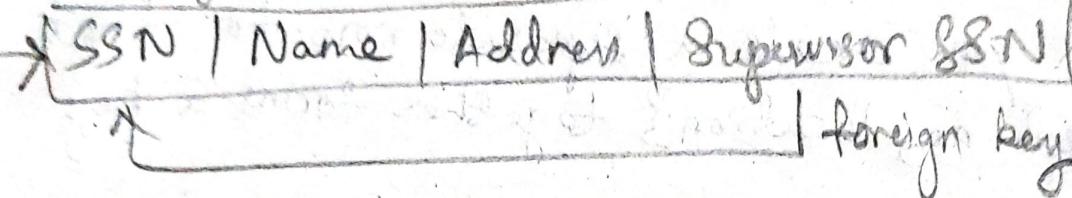
Foreign key



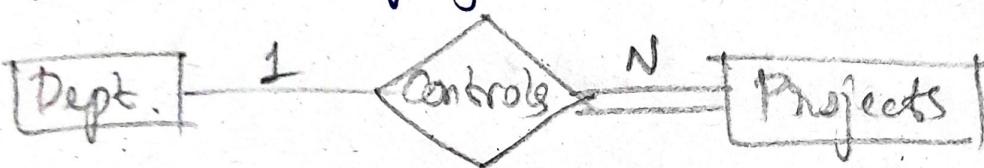
Unary relationship for supervisor +



Employee



Dept. controlled projects +



Dept

D. No.	D. Name	Mgr. SSN

Project

P. No.	P. Name	P. Loc	D. No.

SSN	P. No.	Hours

Foreign Key

Foreign Key

It's always desirable to add an attribute (as a foreign key) to a relation, referencing to another relation, whilst illustrating a relationship.

But in case it's a many-to-many mapping cardinality, we may be forced to show a separate relation illustrating the relationship consisting of collection of primary keys of the referencing to the entities.

Relational Algebra:

DDL commands → Data Definition Language

also called DML → Data Manipulation Language

SQL Schema → Structure of a database

Company (Employee, Dept., Dept. Loc., Project) } together it's
Employee (SSN, Name, Gender...) } called a schema

At different instances, we can have different no of rows(tuples) in our table.

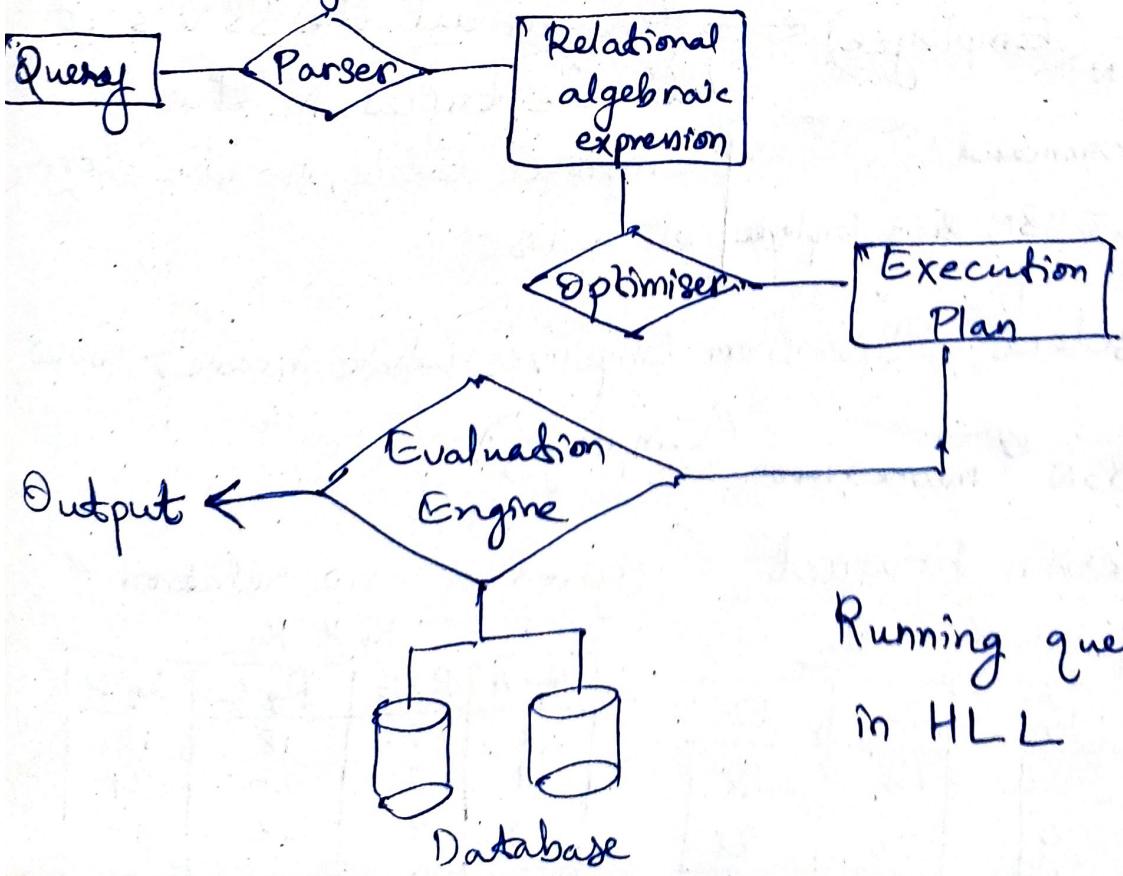
DDL commands → related to constructing the schema of the database

Eg → create, drop & delete table

DML commands → queries to get & insert data, all in all, manipulate the database

Eg → insert & select queries

Relational Algebra → Mathematical equivalent of HLL



Fundamental Operations

1. Select } unary operations
2. Project }
3. Cartesian product
4. Rename
5. Set difference
6. Union

Derived operations

1. Natural join
2. Divide
3. Theta join
4. Assignment
5. Intersection

Partitions table horizontally

Select Operation \equiv where clause in SQL

↳ Unary operation \leftarrow Accepts only 1 relation

Employee: [SSN | Name | Income]

SQL command: select * from Employee where Income > 100

all columns selected \rightarrow all tuples from Employee satisfying

Equivalent relational algebraic expression

$\Sigma_{\text{Income} > 10000} (\text{Employee})$

Project Operation \leftarrow Partitions table vertically

$\Pi_{\text{SSN}} (\text{Employee}) \leftarrow$ Returns all the SSN's of

all the tuples in the Employee table, in the form of a list

SQL command:

select SSN from Employee

\rightarrow select SSN from Employee where income > 10000

$\Pi_{\text{SSN}} \sigma_{\text{Income} > 10000} (\text{Employee})$

Cartesian Product

\rightarrow gives a new relation

R ₁	A B
1	2
3	4

R ₂	C D
18	19
20	21

R ₁ X R ₂		R ₁ .A R ₁ .B	R ₂ .C R ₂ .D
1	2	18	19
1	2	20	21
3	4	18	19
3	4	20	21

We are NOT taking cartesian product of one tuple with one attribute.
It's simply product of one tuple with that of the 2nd relation.

Book(AccNo, Year, Title)

User(CardNo, Name, Address)

Supplier(SName, SAddress)

Borrow(AccNo, CardNo, DOB)

Find the Acc. No. of all books issued by Rakesh Agrawal

Select AccNo from Borrow where CardNo = that of Rakesh

We'll need the Cartesian product (User x Borrow)

Card No	Name
12	Rakesh
13	Ram
14	Shyam

Acc. No.	Card No.	DOB
A101	12	-
A201	12	-
A301	13	-

Select AccNo from Borrow
join User u on b.CardNo = u.CardNo
where u.Name = 'Rakesh'

TI
Borrow. $\sigma_{UserCardNo = b.CardNo}$
AccNo \wedge $u.Name = 'Rakesh'$

User X Borrow

User Card No	User Name	Acc. No.	Card No.
12	Rakesh	A101	12
12	Rakesh	A201	12
12	Rakesh	A301	13

Name = Rakesh

AND

User = Borrow
Card No. = Card No.

and so on.

In SQL, basically join the tables and then do the select operation.

Relational algebraic expressions

$\pi_{\text{Acc. No.}} \sigma_{\text{User Card No.} = \text{Borrow card No.}}$ (\cup User X Borrow)
Name = 'Rakesh'

DBMS

23/8/23

Relational Algebra & Tuple Calculus

↳ Mathematical formulation of commands given in Structured Query Language (SQL)

Cartesian Product involves more than 1 relation, not a unary operation, rather a binary one.

Union: $R_1 \cup R_2$, not $R_1 \times R_2$ SQL command:

union

In order for $R_1 \cup R_2$ to happen, no. of attributes of relation R_1 = no. of attributes of relation R_2

AND corresponding domains of attributes have to be the same.

Degree of a relation = no. of attributes of the relation.

Book(Ace No, Yr of Pub, Title)

User(Card No, Name, Address)

Supplier(S Name, Address) date of

Borrow(Ace NO, Card No, DoI) issue

Supplied By(Ace No, S Name, Price, DoS)

Q. Find all the books which are either issued or have been supplied by a supplier.

All books issued → given by relation Borrow

All books supplied → given by Supplied By

Borrow \cup SuppliedBy can't be done directly, so we just need to take the AccNo attribute, which is what we need. ← done using project operation.

$\Pi_{AccNo} (Borrow) \cup \Pi_{AccNo} (SuppliedBy)$

Find AccNo of all books available in library.

All books whose AccNo is NOT in Borrow.

So, set difference operation.

Book - Borrow, take the equal attributes only -

$\Pi_{AccNo} (Book) - \Pi_{AccNo} (Borrow)$ on the basis of this

Cartesian Product of 2 tables on the basis of some condition + SQL query

Select Book.AccNo from Book, Borrow where

↑
Gives list of AccNo of all borrowed books

Name of users who've borrowed books

$\Pi_{Name} (User)$ (User x Borrow)) Cartesian product done internally
User.CardNo = Borrow.CardNo

Query + Select User.Name from (User, Borrow) where User.CardNo = Borrow.CardNo

For SQL query practice → Korth, Evan Bayross

Rename : P (oldrelation)
newrelation

Eg + P (Book)
Booknew

Renaming attributes as well + new names of attributes old name of relation

P (AccNoNew, YrOfPubNew, TitleNew)(Book)

Booknew ← new name of relation

5/9/23

SQL Operators

`Distinct` keyword → Takes only unique values under columns/s specified, rejects repetitions.

`Count(...)` → Displays no. of concerned elements from column/s specified

State City Select distinct state from table

X A X Select count(distinct state) from table

X B Y ↳ 2

X C

Select distinct city, state from table

Y D

AX Select count(distinct city, state)

Y E

BX from table

CX ↳ 5

DY

EY

Wildcard → '_' → one character

'%' → zero/multiple characters

`Like` operator → for string matching

Comparison operators don't work on null values.

So there you have to use `IS NULL`, `IS NOT NULL`

DBMS

5/9/23

Simple Cartesian product is seldom used because we associate it with several conditions.

Natural Join → It's a combination of project,

select and cartesian product (binary operation).

Say, a natural join of tables R_1 and R_2 would be $R_1 \bowtie R_2$.

Eg:

R_1	
A	B
1	2
3	4
5	6

R_2	
B	C
2	22
4	82

$$\pi_{R_1 \cdot A, R_2 \cdot C}^{\sigma} (R_1 \bowtie R_2) = R_1 \bowtie R_2$$

$R_1 \bowtie R_2$

A	B	C
1	2	22
3	4	82

this meaning is by default implemented by \bowtie , on the basis of common columns

R_1

R_2

A	B	C
1	2	22
3	4	86
5	6	34

B	C
2	22
4	82

As 2 columns are common, instead of similarity on the basis of element in 1 column, now it will be on the basis of an ordered tuple of 2 intersecting (common) columns.

$R_1 \bowtie R_2$

A	B	C
1	2	22

Sailor(sid, sname, Age, Rating)

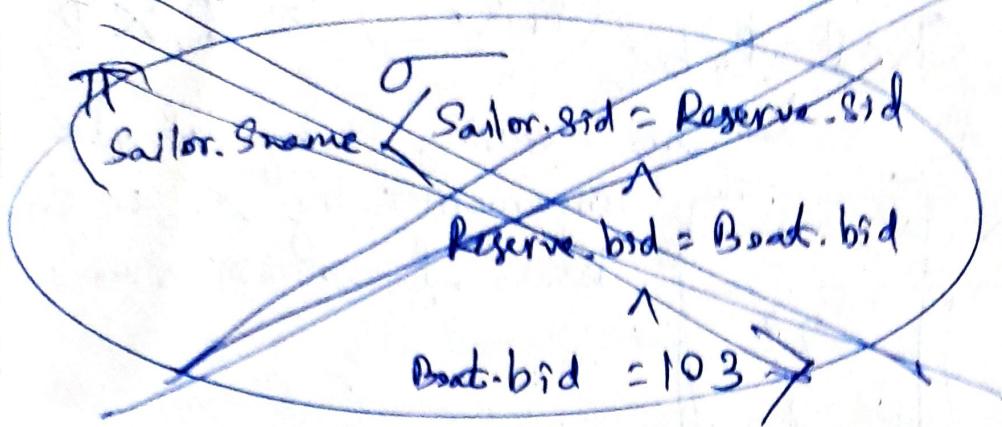
Reserve(sid, bid, day)

Boat(bid, bname, color)

Q1. Find names of sailors who've reserved boat '103'.

Select s. Sname from Sailor s inner join Reserve r on s. Sid = r. Sid inner join Boat b on r. bid = b. bid

where $b.bid = 103$



$\pi_{Sailor.sname} \sigma_{\langle Reserve.bid = 103 \rangle} (Sailor \bowtie Reserve)$

$\pi_{Sname} \sigma_{\langle Sailor.bid = Reserve.sid \rangle} (Sailor \times Reserve)$
 $\sigma_{Reserve.bid = 103}$

~~$\pi_{Sname} ((\sigma_{Reserve.bid = 103}) \bowtie Reserve)$~~

$(Sailor \bowtie Reserve)$ is going to give a much larger table on which you'll perform conditional select.

But this is better as this truncates the Reserve table so natural join product gives a smaller table, hence taking lesser time.
Q2. Find Sailor names who reserved a red boat.

Select $s.sname$ from Sailor & inner join Reserve r on $s.sid = r.sid$ inner join Boat b on $r.bid = b.bid$ where $b.color = 'Red'$

Finding $\sigma_{color = 'Red'}(Boat)$ will be better to optimize.

$\Pi_{\text{Sname}} (\sigma_{\text{color} = \text{'Red'}} (\text{Boat}) \bowtie \text{Reserve}) \bowtie \text{Sailor}$

Select Sname from Sailor s

inner join Reserve r on s.Sid = r.Sid

inner join Boat b on r.bid = b.bid

where b.color = 'Red'

DBMS

6/9/23

Q3. Find names of sailors who've reserved atleast one boat.

Select distinct S-Sname from Sailors inner join Reserve r on S-Sid > r.Sid

$\Pi_{\text{Sname}}^{\text{distinct}} (\text{Sailor} \bowtie \text{Reserve})$

$\Rightarrow \Pi_{\text{Sname}} (\text{Sailor}) - \Pi_{\text{Sname}} (\sigma_{\text{Sname} \neq \text{r.Sid}} (\text{Sailor} \bowtie \text{Reserve}))$

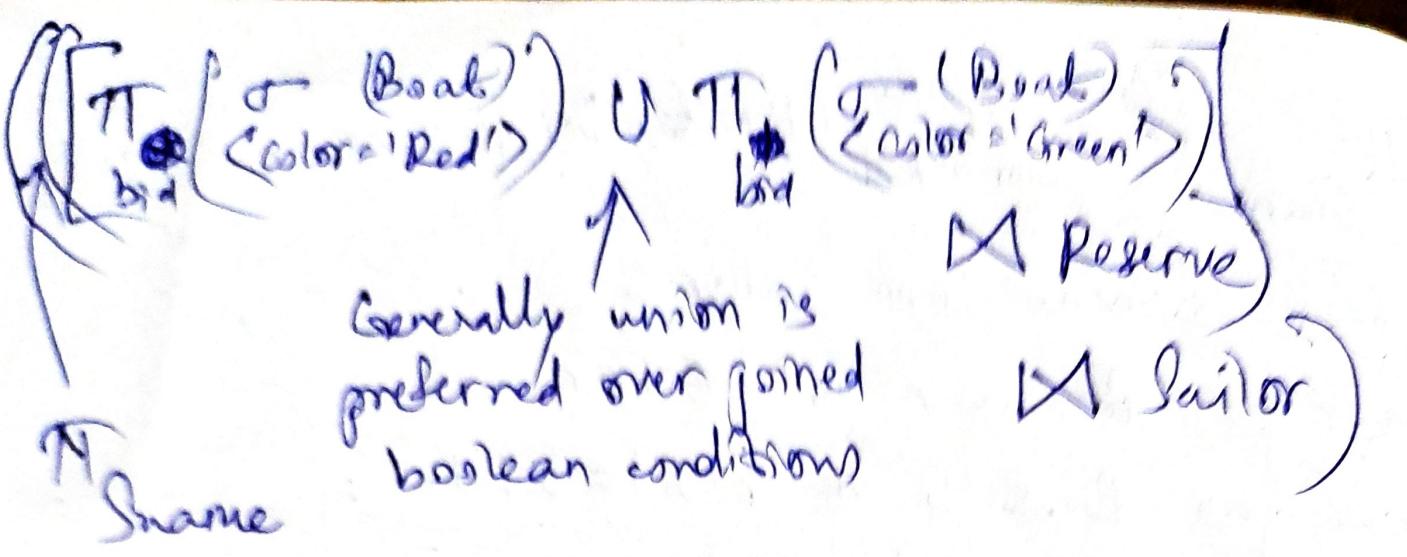
$\nwarrow \text{r.bid} = \text{NULL}$

$\Pi_{\text{Sname}} (\text{Sailor}) - \Pi_{\text{Sname}} (\sigma_{(\text{Sailor}(s) \times \text{Reserve}(r))} (\text{Sailor} \bowtie \text{Reserve}))$

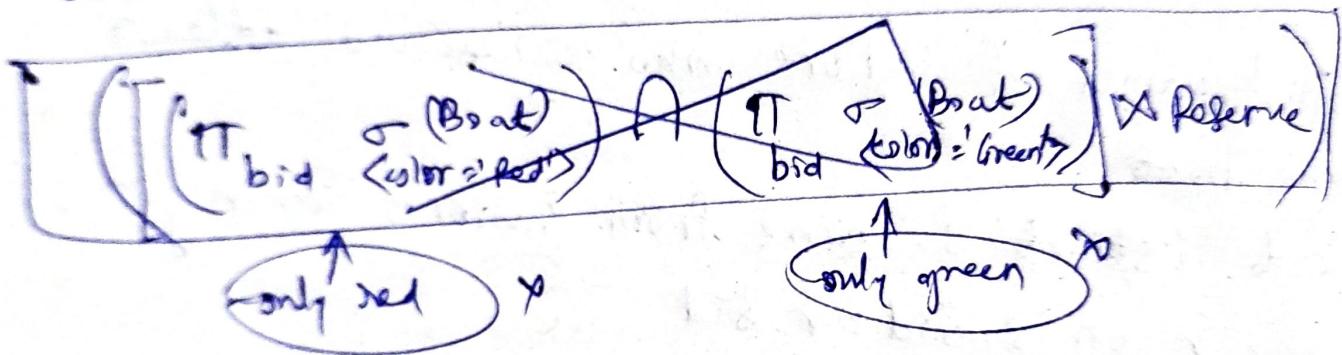
Q4. Find sailor names who've reserved a 'red' boat or green boat

$\Pi_{\text{Sname}} [\sigma_{(\text{Boat})} [\sigma_{\text{color} = \text{'Red'}} (\text{Boat}) \cup \sigma_{\text{color} = \text{'Green'}} (\text{Boat})] \bowtie \text{Reserve} \bowtie \text{Sailor}]$

Columns are unjoined



Q5. Find sailor names who've reserved a Red boat and a Green boat.



Select distinct s-Sname from Sailor s.

inner join (select r.Sid from Reserve r inner join Boats b on r.bid = b.bid where b.color = 'Red' ~~OR~~ b.color = 'Green') \bowtie
 on ~~r.Sid = s.Sid~~

